

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3
        4 import os
        5
        6 # import model related libraries
        7 from sklearn.linear_model import LinearRegression
        8 from sklearn.model_selection import train_test_split
        9
       10 # import module to calculate model performance metrics
       11 from sklearn import metrics
```

```
In [2]: 1 #data_path = "data/Advertising.csv" # or Load the dataset directly from the link
        2 #data_link = "http://www-bcf.usc.edu/~gareth/ISL/Advertising.csv"
        3
        4 #Step 1: get the data (read the csv file)
        5 os.chdir('F:\Training Python')
        6 data = pd.read_csv('Advertising.csv', index_col=0 , error_bad_lines = False)
```

C:\Users\TEMP\AppData\Local\Temp\ipykernel\_7512\2235661692.py:6: FutureWarning: The error\_bad\_lines argument has been deprecated and will be removed in a future version. Use on\_bad\_lines in the future.

```
data = pd.read_csv('Advertising.csv', index_col=0 , error_bad_lines = False)
```

```
In [3]: 1 data
```

```
Out[3]:
```

	TV	Radio	Newspaper	Sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9
...	...	...	...	...
196	38.2	3.7	13.8	7.6
197	94.2	4.9	8.1	9.7
198	177.0	9.3	6.4	12.8
199	283.6	42.0	66.2	25.5
200	232.1	8.6	8.7	13.4

200 rows × 4 columns

## What are the features?

TV: advertising dollars spent on TV for a single product in a given market (in thousands of dollars) Radio: advertising dollars spent on Radio Newspaper: advertising dollars spent on Newspaper What is the response?

Sales: sales of a single product in a given market (in thousands of widgets)

```
In [4]: 1 # print the shape of the DataFrame
        2 data.shape
```

```
Out[4]: (200, 4)
```

```
In [5]: 1 #Step 2: select independent(X) and dependent variable (y)
2 # create a Python list of feature names
3 feature_names=['TV']
4
5 # use the list to select a subset of the original DataFrame
6 X = data[feature_names]
7
8 # sales
9 y = data.Sales
```

```
In [6]: 1 # Step 3: Splitting X and y into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1, test_size=0.2)
```

```
In [7]: 1 # Step 4: Fit linear regression model to training set
2 # Linear Regression Model
3 linreg = LinearRegression()
4
5 # fit the model to the training data (learn the coefficients)
6 linreg.fit(X_train, y_train)
```

Out[7]: LinearRegression()

```
In [8]: 1 print("Intercept=", linreg.intercept_)
2 print("Slope=", linreg.coef_)
```

Intercept= 6.799773449796854  
Slope= [0.0492751]

Interpreting Model Coefficients How do we interpret the TV coefficient ( $\beta_1$ )?

A "unit" increase in TV ad spending is associated with a 0.0492751 "unit" increase in Sales. Or more clearly: An additional \$1,000 spent on TV ads is associated with an increase in sales of 49.2751 widgets. Note that if an increase in TV ad spending was associated with a decrease in sales,  $\beta_1$  would be negative.

Hypothesis Testing and p-values Generally speaking, you start with a null hypothesis and an alternative hypothesis (that is opposite the null). Then, you check whether the data supports rejecting the null hypothesis or failing to reject the null hypothesis.

(Note that "failing to reject" the null is not the same as "accepting" the null hypothesis. The alternative hypothesis may indeed be true, except that you just don't have enough data to show that.)

As it relates to model coefficients, here is the conventional hypothesis test:

null hypothesis: There is no relationship between TV ads and Sales (and thus  $\beta_1$  equals zero) alternative hypothesis: There is a relationship between TV ads and Sales (and thus  $\beta_1$  is not equal to zero) How do we test this hypothesis? Intuitively, we reject the null (and thus believe the alternative) if the 95% confidence interval does not include zero. Conversely, the p-value represents the probability that the coefficient is actually zero:#

```
In [9]: 1 from sklearn.feature_selection import f_regression
2 f_regression=f_regression(X_train, y_train) #returns Fvalues of features; p values of
3 f_regression
```

Out[9]: (array([280.14341559]), array([8.01212052e-37]))

```
In [10]: 1 #Step 5: Test the model's generalization ability using test set
2 # make predictions on the testing set
3 y_pred = linreg.predict(X_test)
```

```
In [11]: 1 y_pred
```

```
Out[11]: array([17.18696362, 16.77798032, 11.51540011, 20.60665526, 19.30579273,
                20.77419058, 14.84639657, 15.70871075, 10.2785952 , 17.41362906,
                14.90552669, 10.20961007, 17.37913649, 12.21017895, 17.92609005,
                12.99365298, 13.28930355, 21.12404376,  8.0612159 , 17.18203611,
                11.74699305, 10.14062494,  8.03657835, 12.09191872, 12.36293175,
                16.08320147,  8.92353007, 19.05941725, 15.01885941, 18.63072392,
                18.6208689 , 18.35478338, 14.17625527, 15.18639473, 19.03970721,
                15.91073864, 17.75855473, 13.17597083, 17.48261419,  7.76556532])
```

```
In [12]: 1 # Step 6: Compute the performance of the model using metrics
        2 print("RMSE=", np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
        3 print("Rsquare=", metrics.r2_score(y_test, y_pred))
```

```
RMSE= 3.2953520791575928
Rsquare= 0.41535307148347866
```

```
In [13]: 1 df_predicted=pd.DataFrame()
        2 df_predicted['Actual']=y_test
        3 df_predicted['Predicted']=y_pred
        4 df_predicted.head(10)
```

```
Out[13]:
```

	Actual	Predicted
59	23.8	17.186964
41	16.6	16.777980
35	9.5	11.515400
103	14.8	20.606655
185	17.6	19.305793
199	25.5	20.774191
96	16.9	14.846397
5	12.9	15.708711
30	10.5	10.278595
169	17.1	17.413629

```
In [ ]: 1
```